

Engram activity and impact

Repository: engram · Report date: 22 September 2026 · Evidence window: 23:20–23:33 UTC, 22 Sep 2026

QUESTION

What does the local record show about Engram's activity and impact in this repository? (No specific question was given, so this report covers the general version.)

Bottom line

Engram is installed and in use here, but the record is short: 18 events over about 13 minutes of one session. It shows a Cortex that is already maintained and one focused task, which Engram estimates cut the Cortex context handed to the agent by 70%. **No paired evaluation has been run, so nothing is known yet about task success, total tokens, cost, or speed.**

Observed activity RECORDED

18	12	1	8
Evidence events	Cortex operations	Focused tasks	Cortex areas mapped
ITEM	WHAT THE RECORD SHOWS		
Setup	Init ran at 23:20 UTC and took 7 ms. It did not create a Cortex because one was already there.		
First maintained Cortex	Not recorded. The Cortex existed before evidence logging began, so its creation falls outside this window. The current Cortex is valid and maps 8 areas.		
Activity by source	CLI 17 · dashboard 1 · MCP 0. So far all agent activity has gone through the CLI; the MCP server has not been used.		

Estimated context reduction ESTIMATE

≈2,237	≈70%
Cortex-context tokens not sent (estimate, all focused tasks)	Average cut in Cortex context per focused task

- This compares the size of the focused Cortex slice with the size of the full Cortex. It is **not** a count of the model's total tokens, and it is not cost, latency, or success.
- It rests on **one** focused task, so treat it as a single example, not an average you can rely on.

Measured task outcomes UNKNOWN

No paired evaluation report exists. The repository defines an evaluation suite (*core-v1*) and adapters, but no results have been recorded. So this record says nothing about whether Engram makes agents more accurate, faster, or cheaper, and this report makes no such comparison.

Limitations

- Very small sample: one session, about 13 minutes, and one focused task.
- The evidence starts after the Cortex was created, so early setup and authoring effort is not captured.
- Estimates only cover the Cortex context Engram supplies. They say nothing about the rest of the agent's prompt or tool use.

Next action: Run the paired *core-v1* evaluation (with Engram vs. without) to measure task success, total tokens, and time. Until then, impact claims should stay limited to the activity and estimates above.

Source: output of engram evidence plus Cortex revision metadata. Raw evidence-log contents, task prompts, and source snippets are left out on purpose.